AI and Model Risk Governance Policy

Policy ID: CRF-RSK-505. Version 2.0. Owner: Model Risk Management. Effective: May 1, 2026.

1. Purpose and Scope

This policy governs the development, validation, deployment and monitoring of models and AI systems at Cobalt Ridge Financial, including generative AI assistants used by employees. It aligns with supervisory guidance on model risk management.

2. Model Inventory and Risk Tiering

2.1 Every model and AI system must be registered in the model inventory before use in production.

2.2 Models are assigned a risk tier from Tier 1 (highest) to Tier 3 (lowest), based on financial impact, customer impact and regulatory exposure.

2.3 Tier 1 models must be independently validated before deployment and revalidated every 12 months. Tier 2 models are revalidated every 24 months and Tier 3 models every 36 months.

3. Generative AI Requirements

3.1 Generative AI assistants must ground answers in approved internal sources and display citations to the source document.

3.2 Personal data must be masked before any prompt is sent to a model, whether hosted internally or by a vendor.

3.3 Generative AI may not make or communicate final credit decisions, adverse action notices, or SAR filing decisions. A qualified human must review and approve these outcomes.

4. Logging and Explainability

4.1 Each generative AI interaction must be logged with the user query, masked query, retrieved sources, model prompt, model response and response latency.

4.2 Interaction logs are retained for 3 years and are available to Internal Audit and Model Risk Management.

5. Monitoring

5.1 Production AI systems are monitored monthly for accuracy, hallucination rate, latency and data leakage.

5.2 A hallucination rate above 2 percent on the monthly evaluation set, or any confirmed leakage of restricted data, triggers an incident review and may require the system to be suspended.

5.3 Vendor-hosted models require a third-party risk assessment before onboarding and an annual review thereafter.

6. Roles

The Model Risk Management team owns this policy. Model owners are accountable for documentation and monitoring. The AI Governance Council approves Tier 1 generative AI use cases.